

Contents

| | |
|--|-----------|
| OH Stats: Complete Learning & Reference Guide | 3 |
| How to Use This Guide | 3 |
| PART I: Statistical Foundations | 4 |
| 1. What Are We Analyzing? | 4 |
| The OH Profile: Multi-Modal Health Data | 4 |
| The Unit of Analysis: Subject x Day | 4 |
| The Two-Package Ecosystem | 4 |
| 2. Why T-Tests Don't Work Here | 4 |
| The Setup | 5 |
| The Independence Problem | 5 |
| What Happens If We Ignore This? | 5 |
| The Coin Flip Analogy | 5 |
| 3. Linear Mixed Models: The Solution | 6 |
| The Key Idea | 6 |
| The Intuitive Model | 6 |
| The Statistical Formula | 6 |
| The ICC: Measuring Clustering | 6 |
| 4. Understanding P-Values and Significance | 7 |
| What a P-Value Actually Means | 7 |
| Common Misinterpretations | 7 |
| Statistical vs. Practical Significance | 8 |
| 5. Effect Sizes: How Big Is the Effect? | 8 |
| Raw Units vs. Standardized | 8 |
| Cohen's d for Mixed Models | 8 |
| Cohen's d Interpretation | 8 |
| 6. The Multiple Testing Problem | 9 |
| The Problem | 9 |
| The Solution: Correction Methods | 9 |
| The Two-Layer Strategy | 9 |
| Critical: Which P-Value Feeds FDR? | 10 |
| 7. Model Assumptions and Diagnostics | 10 |
| The Main Assumptions | 10 |
| Don't Panic About Violations! | 10 |
| Visual Diagnostics (Most Important!) | 10 |
| PART II: The Analysis Workflow | 12 |
| 8. Step 1: Load and Discover Your Data | 12 |
| Load Profiles | 12 |
| Explore Specific Sensors | 12 |
| 9. Step 2: Prepare Your Data | 12 |
| The AnalysisDataset Container | 13 |
| Two Ways to Prepare Data | 13 |
| From Pre-Extracted DataFrame (Maximum Control) | 13 |
| Prepare EMG Data (Convenience Wrapper) | 13 |
| Prepare Any Sensor (Generic) | 14 |

| | |
|---|-----------|
| Prepare Questionnaire Data | 14 |
| Prepare Unified Daily Metrics (Sensors + Workload) | 15 |
| Prepare Single-Instance Metrics (Metadata + IPAQ/OSPAQ) | 15 |
| 10. Step 3: Check Data Quality | 15 |
| The Non-Negotiable Pre-Modeling Checks | 15 |
| What to Look For | 16 |
| 11. Step 4: Fit Models | 16 |
| Single Outcome | 16 |
| Multiple Outcomes (Batch) | 16 |
| Model Options | 17 |
| 12. Step 5: Apply Multiplicity Correction | 17 |
| 13. Step 6: Post-Hoc Contrasts | 17 |
| 14. Step 7: Check Diagnostics | 18 |
| PART III: Interpreting and Reporting Results | 19 |
| 15. Understanding the Output | 19 |
| Coefficients Table | 19 |
| Random Effects | 19 |
| Fit Statistics | 20 |
| 16. Reporting Template | 20 |
| Methods Section | 20 |
| Results Section | 21 |
| What to Report (Checklist) | 21 |
| PART IV: Edge Cases & Troubleshooting | 22 |
| 17. Common Problems and Solutions | 22 |
| 17.1 Missing Days / Unbalanced Data | 22 |
| 17.2 Degenerate Outcomes (No Variance) | 22 |
| 17.3 Convergence Failures | 22 |
| 17.4 EMG Left/Right Correlation (side="both") | 23 |
| 17.5 Skewed Distributions | 23 |
| 17.6 Outliers | 24 |
| 17.7 Likert/Ordinal Data | 24 |
| 17.8 Proportions (0-1 bounded) | 24 |
| 17.9 Small Sample Sizes | 25 |
| 18. Quick Troubleshooting Checklist | 25 |
| PART V: Reference | 26 |
| 19. Data Types and Transform Guide | 26 |
| When to Use Each Transform | 26 |
| Pre-Registered Outcomes | 26 |
| 20. Glossary | 26 |
| 21. Quick Reference Card | 27 |
| Minimal Workflow | 27 |
| Decision Tree | 28 |

OH Stats: Complete Learning & Reference Guide

A comprehensive guide to statistical analysis of Occupational Health data

For beginners learning statistics AND practitioners needing a reference

How to Use This Guide

Navigation Guide

Learning statistics? → Start with Part I, read sequentially

Running an analysis? → Jump to Part II for step-by-step workflow

Troubleshooting? → Part IV has solutions to common problems

Writing a paper? → Part III has reporting templates

PART I: Statistical Foundations

Everything you need to understand WHY we use these methods

1. What Are We Analyzing?

The OH Profile: Multi-Modal Health Data

An OH (Occupational Health) profile contains multiple types of data collected from workers:

| Data Type | Examples | Measurement Scale |
|---------------------------|---------------------------------------|---------------------------------|
| Sensor metrics | EMG, accelerometer, heart rate, noise | Continuous (numeric) |
| Questionnaires | COPSOQ, MUEQ, ROSA, IPAQ, OSPAQ | Ordinal or Continuous |
| Daily self-reports | Workload, pain ratings (NPRS) | Ordinal (Likert 1-5, NPRS 0-10) |
| Environmental | Temperature, humidity | Continuous |

The Unit of Analysis: Subject x Day

💡 Key Concept

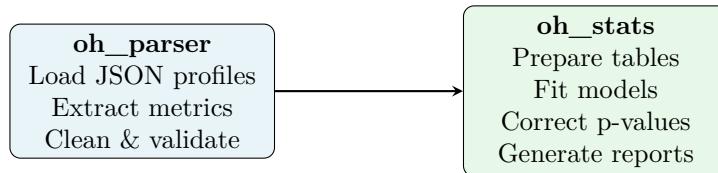
Your primary unit of analysis is **one subject on one day**.

Each row in your analysis = one person's measurements for one day

Key points:

- Each subject contributes **daily aggregated metrics** (e.g., average EMG over the whole day)
- The data are **naturally unbalanced** – some subjects have 3 days, others have 5
- You analyze **each modality separately** (EMG models don't mix with posture or HR models)

The Two-Package Ecosystem



2. Why T-Tests Don't Work Here

The Setup

Imagine you're studying muscle fatigue in office workers. You measure their EMG (muscle activity) every day for a week. Your question: **Does muscle activity change over the week?**

Your data looks like this:

| Subject | Day | EMG_value |
|---------|-----|-----------|
| Alice | 1 | 10.2 |
| Alice | 2 | 9.8 |
| Alice | 3 | 8.5 |
| Bob | 1 | 15.1 |
| Bob | 2 | 14.8 |
| ... | | |

The Independence Problem

⚠ The Problem

Alice's measurements are all related to each other. If Alice naturally has low muscle activity, ALL her measurements will be lower.

T-tests assume **every measurement is independent** – like flipping a coin. But Alice's Day 2 is NOT independent from Alice's Day 1.

Mathematically: The independence assumption fails because $\text{Cov}(Y_{\text{Alice}, \text{Day1}}, Y_{\text{Alice}, \text{Day2}}) \neq 0$.

What Happens If We Ignore This?

| Problem | Effect |
|-------------------------------|---|
| Inflated sample size | We count 320 observations, but really have ~37 independent subjects |
| Too-small p-values | Standard errors underestimated, leading to false confidence |
| False discoveries | We "find" effects that aren't real |
| Unreproducible results | Different samples give wildly different answers |

The Coin Flip Analogy

⚠ Analogy

Imagine you flip a coin 10 times, but you **count each flip 10 times**. You now have "100 observations" but really only 10 independent flips. If you got 6 heads in those 10 real flips, you'd report 60 heads in "100 flips" – and wrongly conclude the coin is biased!

That's exactly what happens when you use t-tests on repeated measures data.

3. Linear Mixed Models: The Solution

The Key Idea

Linear Mixed Models (LMMs) solve this by recognizing **TWO sources of variation**:

$$\text{Total variation} = \text{Between-subject variation} + \text{Within-subject variation}$$

| Component | What it represents |
|------------------------|---|
| Between-subject | Alice vs Bob differences (personal baselines) |
| Within-subject | Day-to-day changes for the same person |

The Intuitive Model

$$\text{EMG_value} = \text{Overall_average} + \text{Day_effect} + \text{Subject's_baseline} + \text{Noise}$$

| Component | What it represents |
|---------------------------|---|
| Overall_average | The typical EMG value across everyone |
| Day_effect | How much Day 2, 3, 4, etc. differ from Day 1 ← <i>this is what we test!</i> |
| Subject's_baseline | Alice is naturally 3 units lower, Bob is 5 units higher, etc. |
| Random_noise | Unexplained day-to-day fluctuations |

The Statistical Formula

For a continuous outcome Y for subject i on day j :

$$Y_{ij} = \beta_0 + \beta_{\text{day}(j)} + u_i + \varepsilon_{ij}$$

Where:

- β_0 = grand mean (intercept)
- $\beta_{\text{day}(j)}$ = effect of day (Day2 vs Day1, Day3 vs Day1, ...)
- $u_i \sim N(0, \sigma_u^2)$ = subject-specific random intercept
- $\varepsilon_{ij} \sim N(0, \sigma^2)$ = residual error

The ICC: Measuring Clustering

The **Intraclass Correlation (ICC)** tells you what proportion of total variation is due to between-subject differences:

$$\text{ICC} = \frac{\sigma_u^2}{\sigma_u^2 + \sigma^2} = \frac{\text{Between-subject variance}}{\text{Total variance}}$$

How to interpret ICC:

| ICC Value | Meaning | Implication |
|-----------|----------------|--|
| 0.0 – 0.2 | Low clustering | Subjects are similar; most variation is day-to-day |
| 0.2 – 0.5 | Moderate | Both sources of variation matter |
| 0.5 – 0.8 | Strong | Who you are matters a lot |
| 0.8 – 1.0 | Very strong | Almost all variation is between people |

❗ Important

In our EMG data, ICC is typically **0.4–0.6**. If ICC is high, you REALLY need mixed models. Using t-tests would give wrong answers.

4. Understanding P-Values and Significance

What a P-Value Actually Means

The **p-value** answers: “*If there were NO real effect, how often would we see data this extreme?*”

| p-value | Interpretation |
|------------|--|
| $p < 0.01$ | Strong evidence of a real effect |
| $p < 0.05$ | Moderate evidence (conventional threshold) |
| $p < 0.10$ | Weak evidence, worth noting but not conclusive |
| $p > 0.10$ | Not enough evidence to claim an effect |

Common Misinterpretations

✖ Common Mistakes

$p < 0.05$ does **NOT** mean “95% sure the effect is real.”

| What people think | Reality |
|--------------------------------------|--------------|
| “95% chance the effect is real” | WRONG |
| “5% chance this is a false positive” | WRONG |

| | |
|--|----------------|
| What people think | Reality |
| “If there’s no effect, we’d see this data only 5% of the time” | CORRECT |

Statistical vs. Practical Significance

A **statistically significant** result ($p < 0.05$) tells you the effect is probably real. It does NOT tell you the effect is **large enough to matter**.

ⓘ Example

“Day 4 is 0.2 %MVC lower than Day 1” might be significant ($p = 0.04$) with large samples.
But is 0.2 %MVC clinically meaningful? That’s a separate question!

Always report effect sizes alongside p-values.

5. Effect Sizes: How Big Is the Effect?

Raw Units vs. Standardized

| Approach | Example | When to use |
|------------------|------------------------------|---|
| Raw units | “Day 4 is 1.93 %MVC lower” | Primary reporting; clinically interpretable |
| Cohen’s d | “ $d = 0.28$ (small effect)” | Cross-study comparison; standardized |

Cohen’s d for Mixed Models

In LMMs, variance is decomposed into components, so there’s no single “pooled SD.” We use **residual-standardized effect size**:

$$d = \frac{\Delta}{\sigma_{\text{residual}}}$$

Where Δ = contrast estimate and σ_{residual} = square root of residual variance.

Cohen’s d Interpretation

| d | |
|-----------|------------|
| < 0.2 | Negligible |
| 0.2 – 0.5 | Small |
| 0.5 – 0.8 | Medium |

| | |
|------------|-------|
| | d |
| ≥ 0.8 | Large |

💡 Recommendation

Always report raw-unit effects with confidence intervals as the primary result. Cohen's d is supplementary for readers who want standardized comparisons across studies.

6. The Multiple Testing Problem

The Problem

You're analyzing 20 different EMG outcomes. Even if NONE of them have real effects, you'll probably find at least one "significant" result just by chance!

Why? If $p < 0.05$ means "5% chance when there's no effect," then:

- Test 1 outcome: 5% chance of false positive
- Test 20 outcomes: $1 - (0.95)^{20} \approx 64\%$ chance of AT LEAST ONE false positive!

The Solution: Correction Methods

| Method | Controls | When to use |
|--------------------------------------|--|-------------------------------------|
| FDR (False Discovery Rate) | Expected proportion of false discoveries | Exploratory analysis, many outcomes |
| FWER (Family-Wise Error Rate) | Chance of ANY false positive | Confirmatory, few primary outcomes |

The Two-Layer Strategy

Our pipeline uses a **two-layer correction**:

Two-Layer Correction Strategy

Layer 1 (across outcomes): FDR on LRT p-values
"Which outcomes show any day effect?"

- ✓ EMG_intensity.mean_percent_mvc: $p_{adj} = 0.005$
- ✓ EMG_apdf.active.p50: $p_{adj} = 0.04$
- ✗ EMG_apdf.rest.p10: $p_{adj} = 0.12$

Layer 2 (within outcome): Holm on post-hoc contrasts
"Which specific days differ?" (only for outcomes that passed Layer 1)

Critical: Which P-Value Feeds FDR?

❗ Critical Distinction

When you see the coefficients table with individual p-values for Day 2, Day 3, etc., we do **NOT** use these for FDR correction.

Instead, we use the **omnibus Likelihood Ratio Test (LRT)** p-value, which asks: "Does including 'day' improve the model AT ALL?"

```
# Access the LRT p-value (this feeds FDR!)
print(result['fit_stats']['lrt_pvalue'])
```

Why the LRT, not coefficient p-values?

- Coefficient p-values test "Day 2 vs Day 1", "Day 3 vs Day 1", etc. – many tests per outcome!
 - LRT asks ONE question per outcome: "Is there ANY day effect?"
 - FDR needs ONE p-value per outcome to work correctly
-

7. Model Assumptions and Diagnostics

The Main Assumptions

| Assumption | What it means | How to check |
|-------------------------------------|--|---------------------------|
| Residuals ~ Normal | "Leftovers" should be bell-shaped | QQ plot, Shapiro-Wilk |
| Constant variance | Spread doesn't change with fitted values | Residuals vs. Fitted plot |
| Independence within clusters | After accounting for subjects, variation is random | Study design |

Don't Panic About Violations!

💡 Good News

LMMs are fairly robust to mild assumption violations.

| Situation | What to do |
|--|--|
| Shapiro-Wilk $p < 0.05$ but QQ plot looks OK | Probably fine, especially with $N > 30$ |
| Moderate skewness ($ skew < 1$) | Usually OK; consider transform if severe |
| A few outliers | Investigate them; run sensitivity analysis |

Visual Diagnostics (Most Important!)

```
import matplotlib.pyplot as plt
from scipy import stats

fig, axes = plt.subplots(1, 2, figsize=(10, 4))

# QQ Plot: Points should follow the diagonal line
stats.probplot(diag['standardized'], dist="norm", plot=axes[0])
axes[0].set_title("QQ Plot (should be a straight line)")

# Residuals vs Fitted: Should be a random cloud around zero
axes[1].scatter(diag['fitted'], diag['residuals'], alpha=0.5)
axes[1].axhline(y=0, color='r', linestyle='--')
axes[1].set_xlabel("Fitted Values")
axes[1].set_ylabel("Residuals")
axes[1].set_title("Residuals vs Fitted (should be random cloud)")

plt.tight_layout()
plt.show()
```

PART II: The Analysis Workflow

Step-by-step guide to running your analysis

8. Step 1: Load and Discover Your Data

Load Profiles

```
from oh_parser import load_profiles
from oh_stats import get_profile_summary, discover_sensors, discover_questionnaires

# Load the OH profiles
profiles = load_profiles("/path/to/OH_profiles")

# FIRST: See what data is available (recommended!)
print(get_profile_summary(profiles))
```

Example output:

OH Profile Summary (42 subjects)

SENSOR DATA:

```
emg: 15 metrics
heart_rate: 8 metrics
noise: 6 metrics
```

SINGLE-INSTANCE QUESTIONNAIRES:

```
personal: 31 fields
biomechanical: 73 fields
```

DAILY QUESTIONNAIRES:

```
workload: 6 fields
pain: 12 fields
```

Explore Specific Sensors

```
# What EMG metrics are available?
sensors = discover_sensors(profiles)
print(sensors['emg'])
# ['EMG_intensity', 'EMG_apdf', 'EMG_muscular_rest', ...]

# What questionnaires?
quests = discover_questionnaires(profiles)
print(quests['single_instance'].keys())
# ['copsoq', 'mueq', 'rosa', 'ipaq', 'ospaq']
```

9. Step 2: Prepare Your Data

The AnalysisDataset Container

AnalysisDataset Structure

All analysis functions expect an **AnalysisDataset** – a standardized container:

- `ds['data']` – The actual DataFrame (long format)
- `ds['outcome_vars']` – List of outcome column names
- `ds['id_var']` – Clustering variable (usually 'subject_id')
- `ds['time_var']` – Time variable (usually 'day_index')
- `ds['grouping_vars']` – Additional grouping (e.g., ["side"])

Two Ways to Prepare Data

Recommended Workflow

Option A: Convenience wrappers (simple cases)

Use `prepare_daily_emg(profiles)` when you want standard extraction.

Option B: From pre-extracted DataFrame (more control)

Use `prepare_from_dataframe(df)` when you've already extracted data with `oh_parser` and want to customize filtering/transformation.

From Pre-Extracted DataFrame (Maximum Control)

If you've already extracted data with `oh_parser`, use `prepare_from_dataframe()`:

```
from oh_parser import extract_nested
from oh_stats import prepare_from_dataframe, fit_lmm

# Step 1: Extract with oh_parser (you control this)
df = extract_nested(
    profiles,
    base_path="sensor_metrics.emg",
    level_names=["date", "level", "side"],
    value_paths=["EMG_intensity.*"],
    flatten_values=True,
)

# Step 2: Apply your custom filtering
df = df[df["level"] == "EMG_daily_metrics"]
df = df.drop(columns=["level"])

# Step 3: Convert to AnalysisDataset (no redundant extraction!)
ds = prepare_from_dataframe(df, sensor="emg", side="average")

# Step 4: Use oh_stats as normal
result = fit_lmm(ds, "EMG_intensity.mean_percent_mvc")
```

Prepare EMG Data (Convenience Wrapper)

```
from oh_stats import prepare_daily_emg

# Keep both sides as separate rows
ds = prepare_daily_emg(profiles, side="both")

# Or average left/right (simpler - RECOMMENDED)
ds = prepare_daily_emg(profiles, side="average")
```

Side handling options:

| Strategy | Effect | When to Use |
|------------------|---------------------------------|--|
| "both" | Left and right as separate rows | When laterality is of interest |
| "average" | Mean of left/right | When laterality is nuisance (recommended) |
| "left" / "right" | Keep only one side | When sides have different meaning |

Prepare Any Sensor (Generic)

```
from oh_stats import prepare_sensor_data

# Heart rate data
hr_ds = prepare_sensor_data(
    profiles,
    sensor="heart_rate",
    base_path="sensor_metrics.heart_rate",
    level_names=["date"],
    value_paths=["HR_BPM_stats.*", "HR_ratio_stats.*"],
)

# Noise data
noise_ds = prepare_sensor_data(
    profiles,
    sensor="noise",
    base_path="sensor_metrics.noise",
    level_names=["date"],
    value_paths=["Noise_statistics.*"],
)
```

Prepare Questionnaire Data

```
from oh_stats import (
    prepare_baseline_questionnaires,
    prepare_daily_pain,
    prepare_daily_workload
)

# Single-instance baseline questionnaires
```

```
baseline_ds = prepare_baseline_questionnaires(profiles, questionnaire_type="copsoq")

# Daily repeated measures
pain_ds = prepare_daily_pain(profiles)
workload_ds = prepare_daily_workload(profiles)
```

Prepare Unified Daily Metrics (Sensors + Workload)

```
from oh_stats import prepare_daily_metrics

# Unified daily dataset with HR, noise, HAR, EMG, and workload
daily_ds = prepare_daily_metrics(profiles)
print(daily_ds["data"].head())
```

Included metrics (when available):

- Workload daily mean (5 fixed items)
- Human activities: sitting/standing/walking durations, sitting proportion, steps
- Heart rate: duration-weighted daily mean/std of HR ratio
- Noise: duration-weighted daily mean/std
- EMG: right-side daily p90/p50 (scaled to 0–1)

HR duration fallback: If watch times are missing, per-day durations fallback to the mean of available HR session durations.

Prepare Single-Instance Metrics (Metadata + IPAQ/OSPAQ)

```
from oh_stats import prepare_single_instance_metrics

# Single-instance metrics (metadata + baseline questionnaires)
single_ds = prepare_single_instance_metrics(profiles)
print(single_ds["data"].head())
```

Included metrics (when available):

- Metadata (all fields under meta_data)
- IPAQ: ordinal level (leve/moderada/alta → 1/2/3) + total_met
- OSPAQ: sitting percentage scaled to 0–1
- Weekly HAR total duration (summed across days)

10. Step 3: Check Data Quality

❗ ALWAYS DO THIS!

Never skip data quality checks before modeling.

The Non-Negotiable Pre-Modeling Checks

```

from oh_stats import summarize_outcomes, check_variance, missingness_report

# 1. Basic summary
summary = summarize_outcomes(ds)
print(summary)

# 2. Check for missing data
missing = missingness_report(ds)
high_missing = missing[missing['pct_missing'] > 10]
if len(high_missing) > 0:
    print(f"[WARNING] High missingness (>10%): {high_missing['outcome'].tolist()}")

# 3. Check for degenerate variables
variance = check_variance(ds)
degenerate = variance[variance['is_degenerate']]['outcome'].tolist()
if degenerate:
    print(f"[EXCLUDE] Cannot model: {degenerate}")

```

What to Look For

| Check | Threshold | Action |
|------------------|----------------------|--|
| Missing data | > 10% | Investigate pattern; is it random or systematic? |
| Degenerate | mode > 95% of values | Exclude from modeling |
| Extreme skewness | skew > 2 | Consider LOG transform |
| Sample size | < 20 subjects | Results may be unstable |

11. Step 4: Fit Models

Single Outcome

```

from oh_stats import fit_lmm

# Fit a Linear Mixed Model
result = fit_lmm(ds, "EMG_intensity.mean_percent_mvc")

# Check convergence
if result['converged']:
    print("Model fitted successfully!")
else:
    print("WARNING: Model had problems converging")
    print(result['warnings'])

```

Multiple Outcomes (Batch)

```
from oh_stats import fit_all_outcomes

# Fit all outcomes
results = fit_all_outcomes(ds, skip_degenerate=True)

# Or limit to specific outcomes
results = fit_all_outcomes(
    ds,
    outcomes=["EMG_intensity.mean_percent_mvc", "EMG_apdf.active.p50"],
    max_outcomes=10
)
```

Model Options

```
# Day as categorical (default) - tests each day vs Day 1
result = fit_lmm(ds, outcome, day_as_categorical=True)

# Day as linear trend - tests linear change per day
result = fit_lmm(ds, outcome, day_as_categorical=False)

# Apply transformation
from oh_stats import TransformType
result = fit_lmm(ds, outcome, transform=TransformType.LOG)

# Exclude side effect
result = fit_lmm(ds, outcome, include_side=False)
```

12. Step 5: Apply Multiplicity Correction

```
from oh_stats import apply_fdr

# Apply FDR correction across outcomes
fdr_results = apply_fdr(results)
print(fdr_results)
```

Output:

| | outcome | p_raw | p_adjusted | significant |
|--------------------------------|---------|--------|------------|-------------|
| EMG_intensity.mean_percent_mvc | | 0.0003 | 0.0015 | True |
| EMG_intensity.max_percent_mvc | | 0.0001 | 0.0015 | True |
| EMG_apdf.active.p10 | | 0.0180 | 0.0360 | True |
| EMG_apdf.active.p50 | | 0.0712 | 0.0712 | False |

13. Step 6: Post-Hoc Contrasts

⚠ Important

Only run post-hocs for outcomes that passed FDR correction!

```
from oh_stats import pairwise_contrasts

# Get specific day comparisons
contrasts = pairwise_contrasts(result, "day_index", ds, adjustment="holm")
print(contrasts[["contrast", "estimate", "p_adjusted", "cohens_d"]])
```

Output:

| | contrast | estimate | p_adjusted | cohens_d |
|---|-----------|----------|------------|-------------------------|
| 0 | Day1-Day2 | -0.411 | 0.618 | -0.059 |
| 1 | Day1-Day3 | -0.028 | 0.973 | -0.004 |
| 2 | Day1-Day4 | -1.931 | 0.043 | -0.276 <-- Significant! |
| 3 | Day1-Day5 | -1.643 | 0.184 | -0.235 |

14. Step 7: Check Diagnostics

```
from oh_stats import residual_diagnostic

diag = residual_diagnostic(result)

print(f"Normality test p-value: {diag['normality_p']:.4f}")
print(f"Number of outliers: {diag['n_outliers']}")
print(f"Assumptions broadly met: {diag['assumptions_met']}
```

PART III: Interpreting and Reporting Results

15. Understanding the Output

Coefficients Table

| term | estimate | std_error | z_value | p_value | ci_lower | ci_upper |
|-------------------|----------|-----------|---------|---------|----------|----------|
| Intercept | 9.406 | 1.035 | 9.087 | 0.000 | 7.377 | 11.434 |
| C(day_index)[T.2] | -0.411 | 0.825 | -0.498 | 0.618 | -2.029 | 1.206 |
| C(day_index)[T.3] | -0.028 | 0.839 | -0.033 | 0.973 | -1.672 | 1.616 |
| C(day_index)[T.4] | -1.931 | 0.840 | -2.298 | 0.022 | -3.577 | -0.284 |
| C(day_index)[T.5] | -1.643 | 0.975 | -1.685 | 0.092 | -3.554 | 0.268 |
| C(side)[T.right] | 0.902 | 0.550 | 1.641 | 0.101 | -0.175 | 1.980 |

How to read this:

| Column | What it means |
|----------------|---|
| term | What's being compared |
| estimate | The size of the difference (in raw units) |
| std_error | How uncertain we are (smaller = more confident) |
| z_value | Test statistic (estimate / std_error) |
| p_value | Probability this is just random chance |
| ci_lower/upper | 95% confidence interval |

Interpreting each row:

| Row | Interpretation |
|----------------------------|---|
| Intercept (9.406) | Mean %MVC on Day 1, Left side |
| C(day_index)[T.2] = -0.411 | Day 2 is 0.41 units LOWER than Day 1 (not significant) |
| C(day_index)[T.4] = -1.931 | Day 4 is 1.93 units LOWER than Day 1 (p=0.02, significant!) |
| C(side)[T.right] = 0.902 | Right side is 0.90 units HIGHER than left (not significant) |

Random Effects

```
print(result['random_effects'])
# {'group_var': 24.05, 'residual_var': 23.88, 'icc': 0.502}
```

| Component | Value | Meaning |
|--------------|-------|---|
| group_var | 24.05 | Between-subject variance (σ_u^2) |
| residual_var | 23.88 | Within-subject variance (σ^2) |
| icc | 0.502 | 50% of variation is between subjects |

✓ Validation

ICC of 0.50 tells us: Mixed models were definitely the right choice! Half of all variation is just "who the person is."

Fit Statistics

```
print(result['fit_stats'])
# {'aic': 478.4, 'bic': 502.1, 'loglik': -234.2,
#  'lrt_stat': 12.5, 'lrt_df': 4, 'lrt_pvalue': 0.014}
```

| Statistic | Use |
|------------|--|
| AIC/BIC | Compare models (lower = better) |
| loglik | Log-likelihood (for advanced comparisons) |
| lrt_pvalue | The p-value used for FDR correction |

16. Reporting Template

Methods Section

📄 Example Methods Text

Daily EMG metrics were analyzed using linear mixed models with day as a fixed effect (categorical) and random intercepts for subjects to account for repeated measurements within individuals. Side (left/right) was included as a fixed effect. Models were fitted using maximum likelihood estimation via statsmodels (Python). Given the exploratory nature of the analysis across N=10 EMG outcomes, p-values were adjusted using the Benjamini-Hochberg procedure to control the false discovery rate at 5%. Post-hoc pairwise comparisons between days were corrected using the Holm method. Effect sizes were calculated as Cohen's d using the residual standard deviation as the denominator.

Results Section

Example Results Text

We analyzed 320 observations from 37 subjects over 5 monitoring days (mean 4.3 days per subject, range 3–5). The intraclass correlation was 0.50 (95% CI: 0.35–0.65), indicating that 50% of the variance in EMG intensity was attributable to between-subject differences, justifying the use of mixed models.

After FDR correction, 4 of 10 EMG outcomes showed significant day effects (all $p_{\text{adj}} < 0.05$). For mean %MVC specifically, the overall day effect was significant (LRT $\chi^2(4) = 12.5$, $p = 0.014$). Post-hoc comparisons (Holm-adjusted) revealed that Day 4 was significantly lower than Day 1 ($\Delta = -1.93$ %MVC, 95% CI: -3.58 to -0.28, Cohen's $d = 0.28$, $p_{\text{adj}} = 0.043$), representing a small effect.

What to Report (Checklist)

| Element | Example | Where |
|------------------|---|-----------------|
| Sample size | “37 subjects, 320 observations” | Methods/Results |
| ICC | “ICC = 0.50” | Results |
| FDR method | “Benjamini-Hochberg” | Methods |
| Omnibus test | “LRT $\chi^2(4) = 12.5$, $p = 0.014$ ” | Results |
| Effect estimate | “ $\Delta = -1.93$ %MVC” | Results |
| 95% CI | “95% CI: -3.58 to -0.28” | Results |
| Effect size | “Cohen's $d = 0.28$ ” | Results |
| Adjusted p-value | “ $p_{\text{adj}} = 0.043$ ” | Results |

PART IV: Edge Cases & Troubleshooting

17. Common Problems and Solutions

17.1 Missing Days / Unbalanced Data

The situation: Some subjects have 5 days of data, others have only 3.

✓ Good News

LMMs handle this naturally! They use all available data and don't require balanced designs.

What to watch for:

- Is missingness random or systematic? (e.g., do subjects drop out because they're injured?)
- Very few observations per subject (< 3) may cause convergence issues

```
# Check missingness patterns
missing = missingness_report(ds)
print(missing[missing['pct_missing'] > 10])
```

17.2 Degenerate Outcomes (No Variance)

The situation: An outcome is nearly constant (e.g., 95% of values are zero).

The problem: No variance = nothing to model.

Solution: Exclude these outcomes from analysis.

```
variance = check_variance(ds)
degenerate = variance[variance['is_degenerate']]
print(f"Exclude: {list(degenerate['outcome'])}")
```

17.3 Convergence Failures

The situation: `result['converged'] = False`

❗ Warning

The optimizer couldn't find a stable solution. Results are unreliable.

What to try:

1. **Simplify the model:** Remove interactions, use `side="average"`
2. **Check for degenerate outcomes:** Near-constant values cause problems
3. **Check sample size:** Need enough subjects (ideally 20+)
4. **Look at warnings:** `result['warnings']` often explains the issue

```

if not result['converged']:
    print("Warnings:", result.get('warnings', []))
    # Try simpler model
    ds_simple = prepare_daily_emg(profiles, side="average")
    result = fit_lmm(ds_simple, outcome)

```

17.4 EMG Left/Right Correlation (side="both")

The situation: You kept both sides as separate rows, but left and right from the same subject-day are correlated.

The problem: A subject-only random intercept doesn't fully capture same-day correlations.

Three defensible strategies:

| Strategy | Pros | Cons |
|--------------------------|--------------------------------|---------------------------------|
| side="average" | Simplest, no correlation issue | Loses side-specific information |
| Analyze sides separately | Clean interpretation | Doubles the number of tests |
| Keep side="both" | More power | Slight model misspecification |

💡 Recommendation

Start with `side="average"` for simplicity.

17.5 Skewed Distributions

The situation: Residuals are not normally distributed (e.g., right-skewed EMG data).

Don't panic! LMMs are fairly robust to moderate non-normality, especially with larger samples.

When to act:

- Severe skewness (> 2) with small samples
- Heavy ceiling/floor effects

Solutions:

```

import numpy as np

# Log transform (for positive values, especially right-skewed)
ds['data']['log_outcome'] = np.log1p(ds['data'][outcome])

# Or specify in fit_lmm
from oh_stats import TransformType
result = fit_lmm(ds, outcome, transform=TransformType.LOG1P)

```

17.6 Outliers

The situation: A few extreme values are pulling the model.

How to identify:

```
diag = residual_diagnostics(result)
print(f"Outliers (|z| > 3): {diag['n_outliers']}")

# See which observations
import numpy as np
outlier_idx = np.abs(diag['standardized']) > 3
print(ds['data'][outlier_idx])
```

What to do:

1. **Investigate:** Are they data errors or real extreme values?
2. **Sensitivity analysis:** Run with and without outliers
3. **Report both:** “Results were similar with outliers excluded (N=2)”

17.7 Likert/Ordinal Data

The situation: You have questionnaire items on a 1-5 or 0-10 scale.

The theoretical issue: Likert scales are ordinal, not continuous. The difference between 1\$→2 isn't necessarily the same as 4\$→5.

Practical guidance:

| Distribution | Recommendation |
|--------------------------------------|--|
| Roughly symmetric, no ceiling/floor | Treat as continuous with LMM (common practice) |
| Heavy ceiling (most responses = max) | Consider ordinal models or dichotomize |
| Heavy floor (most responses = min) | Consider ordinal models or dichotomize |

If treating as continuous: Always report medians and IQR alongside means.

17.8 Proportions (0-1 bounded)

The situation: Your outcome is a proportion (e.g., % time in a posture).

The problem: Values bounded at 0 and 1; residuals can't be normal at the extremes.

Solution: LOGIT transform

```
# Automatic via registry for registered proportions
result = fit_lmm(ds, 'ospaq_sitting_pct') # Auto-applies LOGIT

# Or manual
result = fit_lmm(ds, outcome, transform=TransformType.LOGIT)
```

17.9 Small Sample Sizes

The situation: You have fewer than 30 subjects.

The problem: Random effect variance estimates become imprecise; model may not converge.

```
if result['n_groups'] < 30:  
    print(f"Warning: Only {result['n_groups']} subjects")  
    print("Consider: wider CIs, simpler models, sensitivity analyses")
```

Recommendations:

- Prefer simpler models (fewer fixed effects)
- Consider reporting alongside bootstrap CIs
- Be cautious about random effect variance interpretations

18. Quick Troubleshooting Checklist

Troubleshooting Checklist

- Model didn't converge?**
→ Try side="average", check for degenerate outcomes, simplify model
- Residuals look weird?**
→ Check for outliers, consider transformation
- Unexpected results?**
→ Check missingness patterns, verify data quality
- p-values all non-significant but you expected effects?**
→ Check ICC (high ICC = less power), check sample size
- Too many significant results?**
→ Are you using FDR correction? Check for data leakage

PART V: Reference

19. Data Types and Transform Guide

When to Use Each Transform

| Outcome Type | Transform | When to Use |
|--------------------------------|--------------|--|
| Continuous (unbounded) | NONE | Default for %MVC, BPM, etc. |
| Right-skewed continuous | LOG or LOG1P | When distribution has long right tail |
| Proportions (0-1) | LOGIT | % time, rest_percent, OSPAQ |
| Counts | LOG1P | Number of events (pragmatic fallback) |
| Ordinal (5+ levels) | NONE | NPRS 0-10, ROSA 1-10 (treat as continuous) |

Pre-Registered Outcomes

| Outcome | Type | Transform |
|------------------------------------|------------|-----------|
| EMG_intensity.mean_percent_mvc | CONTINUOUS | NONE |
| EMG_intensity.iemg_percent_seconds | CONTINUOUS | LOG |
| EMG_apdf.rest_percent | PROPORTION | LOGIT |
| EMG_muscular_rest.gap_count | COUNT | LOG1P |
| copsoq_* | CONTINUOUS | NONE |
| mueq_* | CONTINUOUS | NONE |
| rosa_total | ORDINAL | NONE |
| ipaq_met_min_week | CONTINUOUS | LOG1P |
| ospaq_*_pct | PROPORTION | LOGIT |
| nprs_* | ORDINAL | NONE |

20. Glossary

| Term | Definition |
|------------------------|---|
| AIC | Akaike Information Criterion. Lower = better model fit. |
| AnalysisDataset | Standardized container for analysis-ready data. |

| Term | Definition |
|---------------------------------|--|
| Coefficient | Estimated size of an effect (e.g., Day 4 is -1.93 lower than Day 1). |
| Cohen's d | Standardized effect size: difference / standard deviation. |
| Confidence Interval (CI) | Range that probably contains the true effect. 95% CI means 95% confident. |
| Converged | Model successfully found a solution. If FALSE, results unreliable. |
| FDR | False Discovery Rate. Controls expected proportion of false positives. |
| Fixed Effect | Something we're interested in measuring (e.g., day effect, side effect). |
| FWER | Family-Wise Error Rate. Controls chance of ANY false positive. |
| ICC | Intraclass Correlation. Proportion of variance due to between-subject differences. |
| LMM | Linear Mixed Model. Handles repeated measures via fixed + random effects. |
| LRT | Likelihood Ratio Test. Compares nested models to test if a factor matters. |
| p-value | Probability of seeing your data if there were no real effect. |
| Random Effect | Variation we account for but don't directly measure (e.g., subject baselines). |
| Residual | The “leftover” after the model’s prediction. |
| Transform | Converting data (e.g., LOG) to make it better behaved for modeling. |

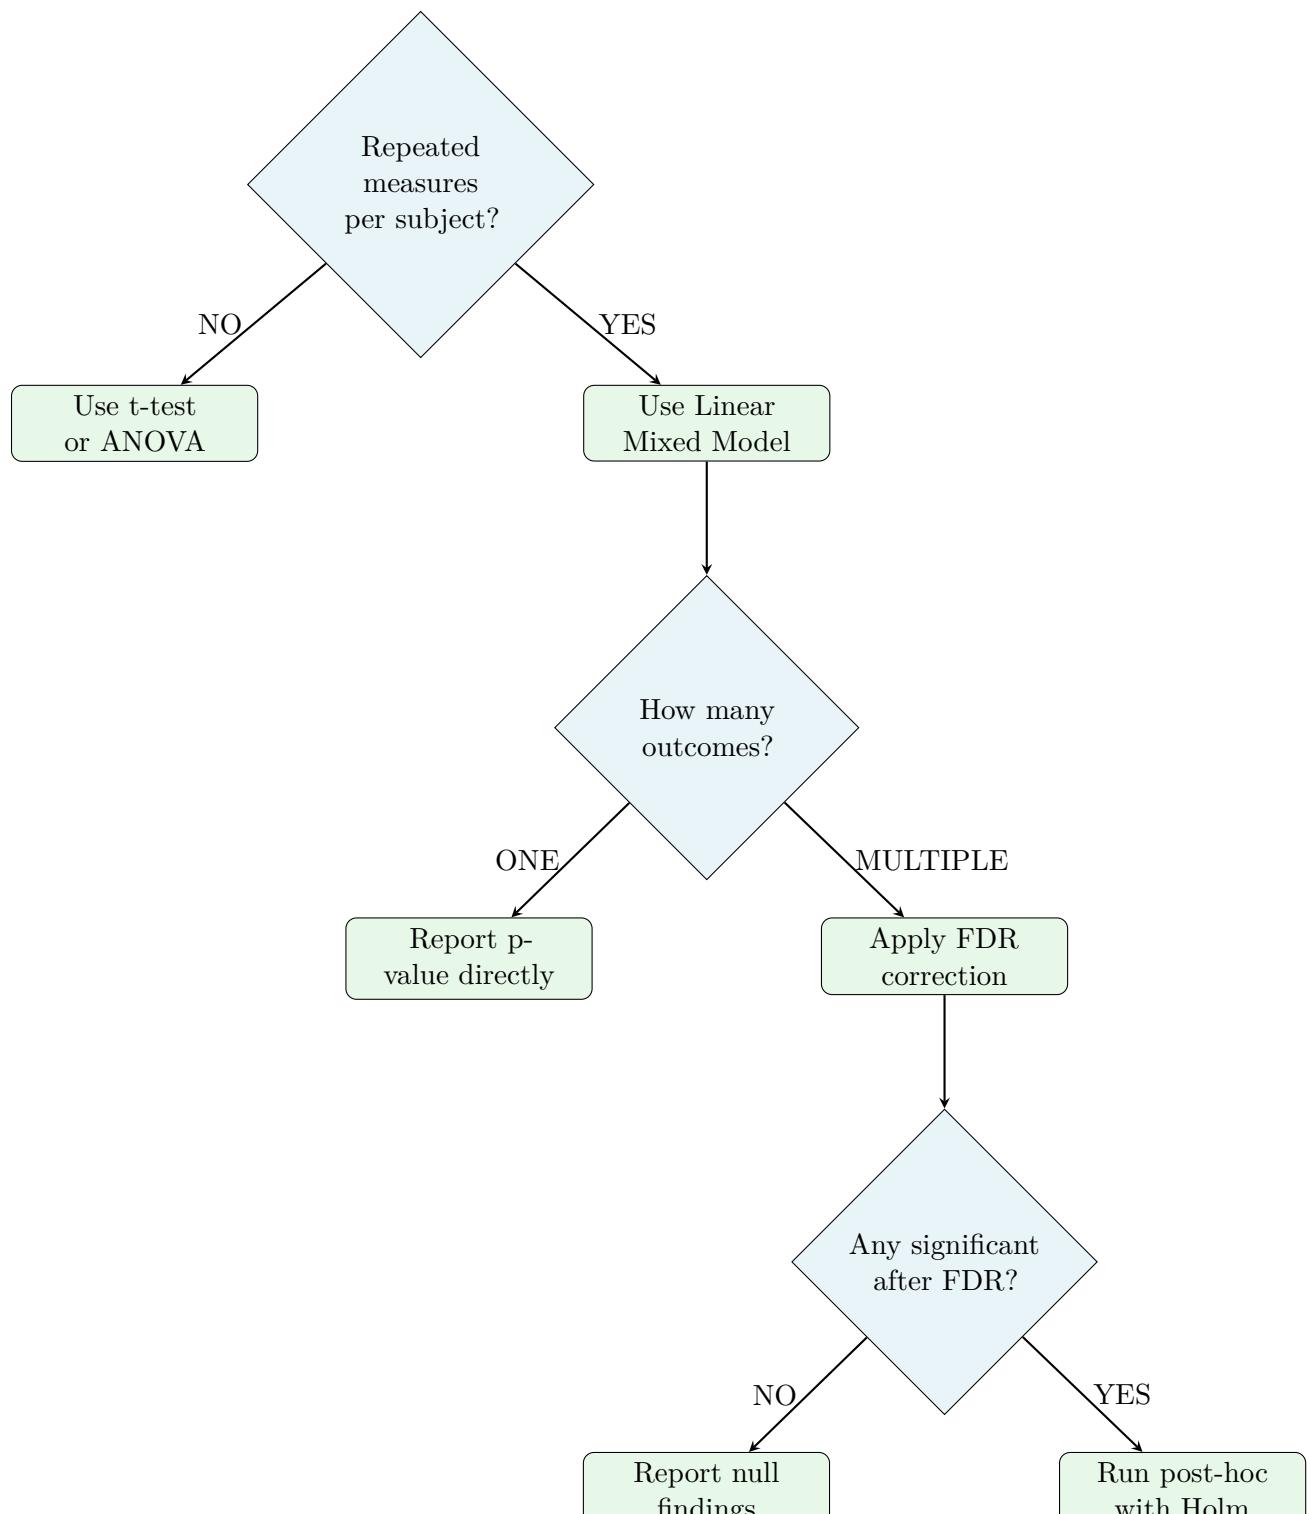
21. Quick Reference Card

Minimal Workflow

```
from oh_parser import load_profiles
from oh_stats import (
    get_profile_summary,
    prepare_daily_emg,
    prepare_daily_metrics,
    prepare_single_instance_metrics,
    summarize_outcomes,
    check_variance,
    fit_all_outcomes,
    apply_fdr,
```

```
)  
  
# 1. Load & Discover  
profiles = load_profiles("/path/to/data")  
print(get_profile_summary(profiles))  
  
# 2. Prepare  
ds = prepare_daily_emg(profiles, side="average")  
  
# Or unified daily metrics  
daily_ds = prepare_daily_metrics(profiles)  
  
# Single-instance metrics  
single_ds = prepare_single_instance_metrics(profiles)  
  
# 3. Check Quality  
print(summarize_outcomes(ds))  
print(check_variance(ds))  
  
# 4. Model  
results = fit_all_outcomes(ds, skip_degenerate=True)  
  
# 5. Correct  
fdr = apply_fdr(results)  
  
# 6. Report  
print(fdr[fdr['significant']])
```

Decision Tree



OH Stats Complete Guide v1.0

January 2026

For oh_stats package v0.3.0